

# Approximate Inference for Neural Networks

Mauricio A. Álvarez PhD,  
H.F. Garcia   C. Guarnizo (TA)



Universidad Tecnológica de Pereira, Pereira, Colombia

## 1 EM algorithm

## 2 Variational EM

## 3 Approximate inference for Neural Networks

- Bayesian Deep Neural Networks
- Laplace Approximation



## 1 EM algorithm

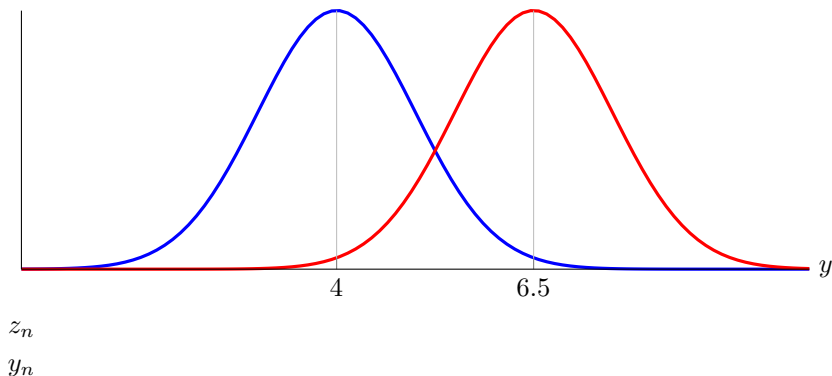
## 2 Variational EM

## 3 Approximate inference for Neural Networks

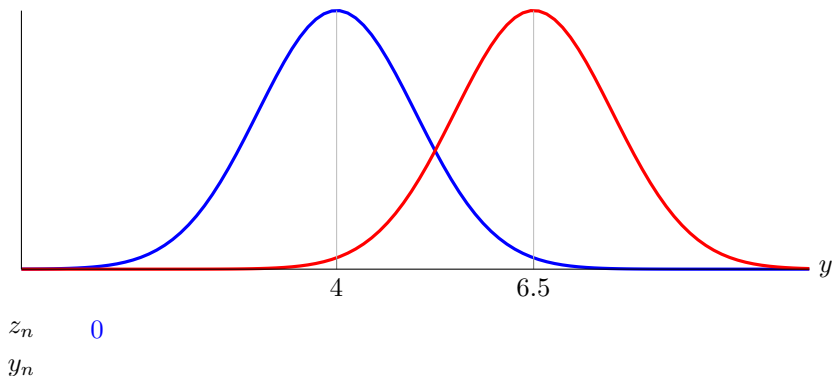
- Bayesian Deep Neural Networks
- Laplace Approximation



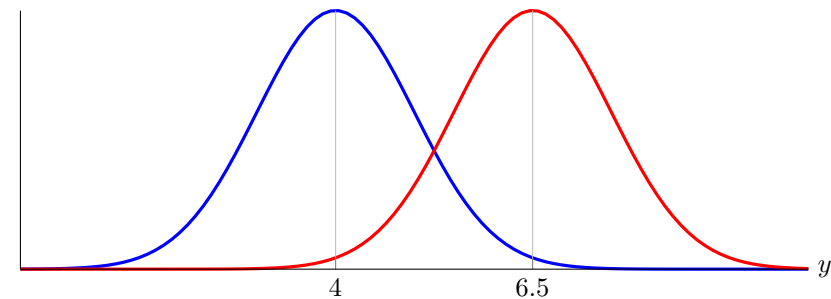
# Problem definition - Example



# Problem definition - Example



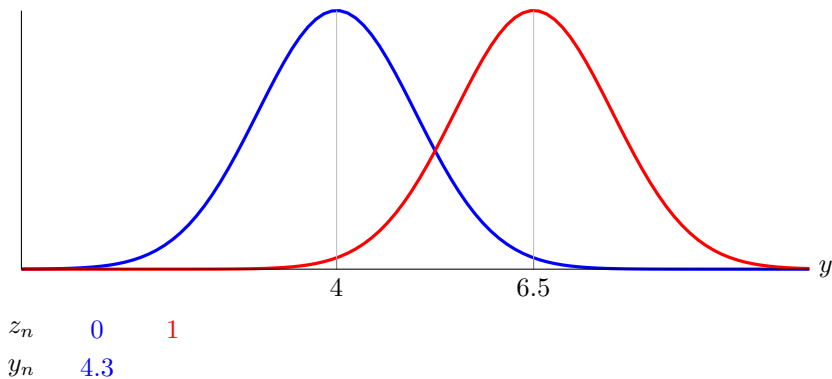
# Problem definition - Example



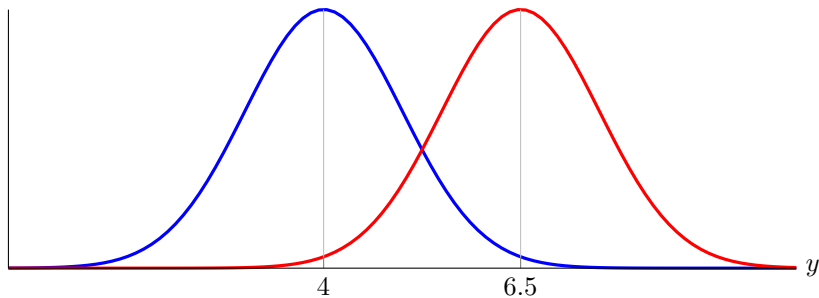
$z_n$      0  
 $y_n$      4.3



# Problem definition - Example



# Problem definition - Example

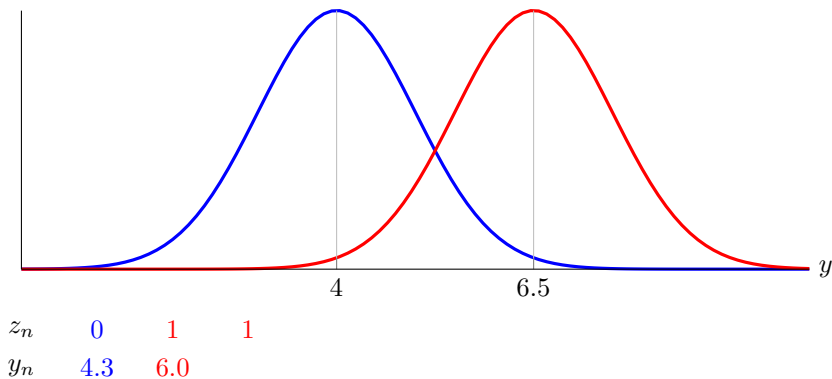


$z_n$	0	1
$y_n$	4.3	6.0

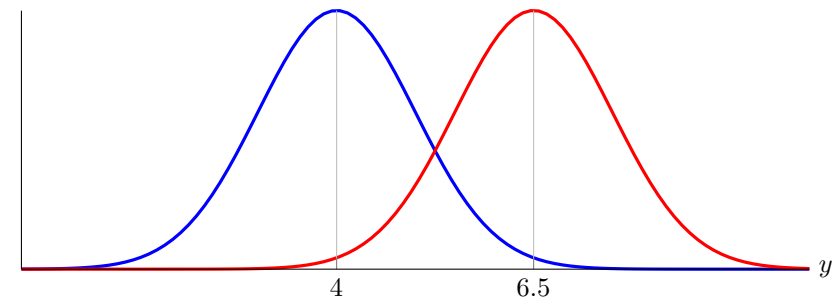




# Problem definition - Example



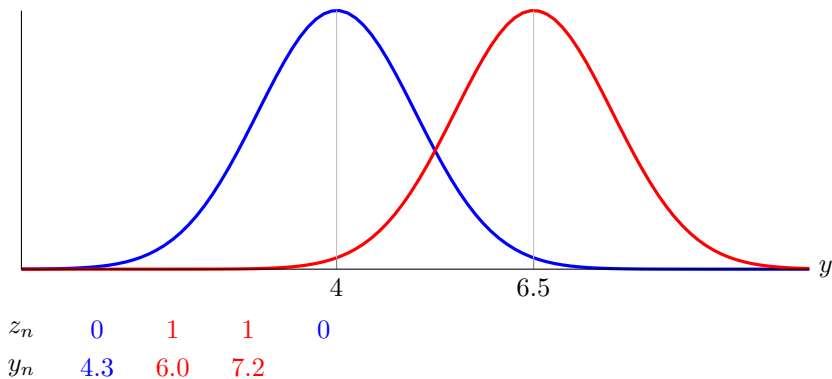
# Problem definition - Example



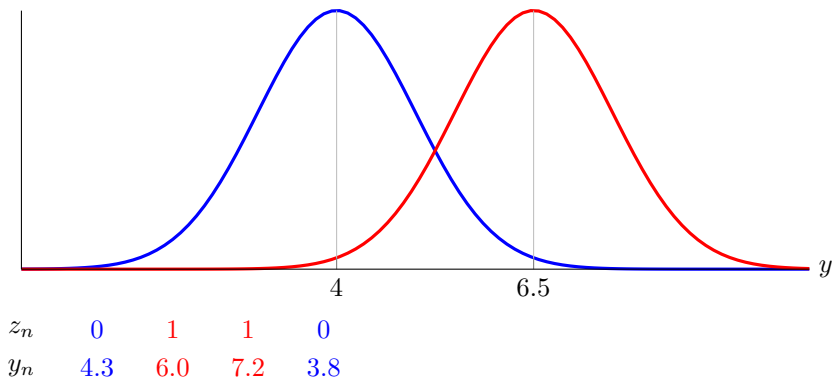
$z_n$	0	1	1
$y_n$	4.3	6.0	7.2



# Problem definition - Example



# Problem definition - Example



# Problem definition - Example

For  $z$ :

$$p(z) = \pi^z (1 - \pi)^{1-z}, \quad p(z = 1) = \pi, \quad p(z = 0) = 1 - \pi.$$

For  $y$ :

$$p(y|z) = (\mathcal{N}(y|\mu_1, \sigma_1^2))^z (\mathcal{N}(y|\mu_2, \sigma_2^2))^{1-z}.$$

The the joint probability:

$$p(y, z) = (\pi \mathcal{N}(y|\mu_1, \sigma_1^2))^z ((1 - \pi) \mathcal{N}(y|\mu_2, \sigma_2^2))^{1-z}.$$

Additionally:

$$\begin{aligned} p(z = 1|y) &= \frac{p(y|z = 1)p(z = 1)}{p(y)} \\ &= \frac{\mathcal{N}(y|\mu_1, \sigma_1^2)\pi}{\pi \mathcal{N}(y|\mu_1, \sigma_1^2) + (1 - \pi) \mathcal{N}(y|\mu_2, \sigma_2^2)} \end{aligned}$$



# Problem definition - Example

Let's assume that we have  $(\mathbf{y}, \mathbf{z}) = \{y_n, z_n\}_{n=1}^N$ .

The parameters  $\theta = \{\pi, \mu_1, \mu_2, \sigma_1^2, \sigma_2^2\}$ , can be estimated as

$$\begin{aligned}\theta_{\text{ML}} &= \arg \max_{\theta} \ln (p(\mathbf{y}, \mathbf{z}|\theta)), \\ &= \arg \max_{\theta} \ln \left( \prod_{n=1}^N p(y_n, z_n|\theta) \right), \\ &= \arg \max_{\theta} \sum_{n=1}^N \ln (p(y_n, z_n|\theta)).\end{aligned}$$

Where:

$$\begin{aligned}\ln p(y_n, z_n|\theta) &= z_n \left( \ln(\pi) + \ln \mathcal{N}(y_n|\mu_1, \sigma_1^2) \right) \\ &\quad + (1 - z_n) \left( \ln(1 - \pi) + \ln \mathcal{N}(y_n|\mu_2, \sigma_2^2) \right).\end{aligned}$$



# Problem definition - Example

Let's assume that  $\mathbf{y} = \{y_n\}_{n=1}^N$ ,  $\mathbf{y}$   $\mathbf{z}$  is unknown (latent).  
The parameters  $\theta = \{\pi, \mu_1, \mu_2, \sigma_1^2, \sigma_2^2\}$ , can be estimated as

$$\begin{aligned}\theta_{\text{ML}} &= \arg \max_{\theta} \sum_{n=1}^N \ln (p(y_n|\theta)) , \\ &= \arg \max_{\theta} \sum_{n=1}^N \ln \left( \sum_z p(y_n, z_n|\theta) \right) .\end{aligned}$$

Where:

$$\ln \left( \sum_z p(y_n, z_n|\theta) \right) = \ln \left( \pi \mathcal{N}(y_n|\mu_1, \sigma_1^2) + (1 - \pi) \mathcal{N}(y_n|\mu_2, \sigma_2^2) \right) .$$



# Problem definition - General

The model consists of observations  $\mathbf{y}$  and a latent random variable  $\mathbf{z}$ . Then

$$\begin{aligned} p(\mathbf{y}|\theta) &= \prod_{n=1}^N p(y_n|\theta) \\ &= \prod_{n=1}^N \sum_{\mathbf{z}} p(y_n, \mathbf{z}|\theta) = \prod_{n=1}^N \sum_{\mathbf{z}} p(y_n|\mathbf{z}, \theta) p(\mathbf{z}|\theta) \end{aligned}$$

The estimation of  $\theta_{\text{ML}}$  is given by

$$\begin{aligned} \theta_{\text{ML}} &= \arg \max_{\theta} \ln (p(\mathbf{y}|\theta)) \\ &= \arg \max_{\theta} \sum_{n=1}^N \ln \left( \sum_{\mathbf{z}} p(y_n|\mathbf{z}, \theta) p(\mathbf{z}|\theta) \right) \end{aligned}$$





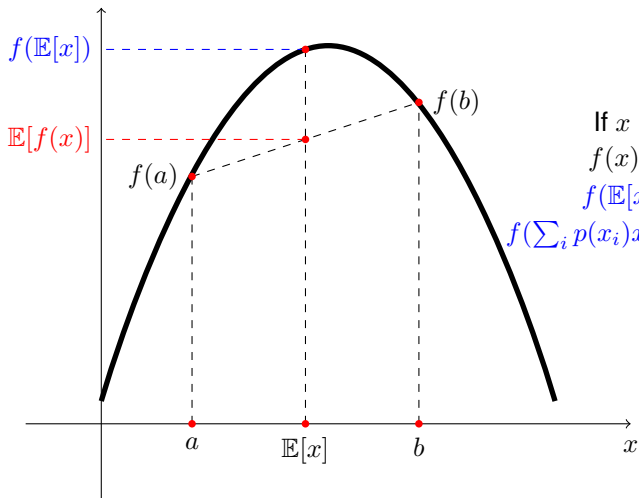
# Problem definition

$$\theta_{\text{ML}} = \arg \max_{\theta} \sum_{n=1}^N \ln \left( \sum_{\mathbf{z}} p(y_n, |\mathbf{z}, \theta) p(\mathbf{z}|\theta) \right)$$

- The sum over  $\mathbf{z}$  couples the parameters  $\theta$ .
- Gradients have no closed form.
- $\mathbf{z}$  and  $\theta$  are also coupled.



# Jensen's Inequality



If  $x$  is a r.v. and  
 $f(x)$  is concave:

$$f(\mathbb{E}[x]) \geq \mathbb{E}[f(x)]$$

$$f(\sum_i p(x_i)x_i) \geq \sum_i p(x_i)f(x_i)$$

# EM algorithm - Jensen's Inequality

$$\log p(\mathbf{y}|\theta) = \sum_{n=1}^N \ln \left( \sum_{\mathbf{z}} p(y_n, \mathbf{z}|\theta) \right)$$



# EM algorithm - Jensen's Inequality

$$\begin{aligned}\log p(\mathbf{y}|\theta) &= \sum_{n=1}^N \ln \left( \sum_{\mathbf{z}} p(y_n, \mathbf{z}|\theta) \right) \\ &= \sum_{n=1}^N \ln \left( \sum_{\mathbf{z}} q(\mathbf{z}) \frac{p(y_n, \mathbf{z}|\theta)}{q(\mathbf{z})} \right)\end{aligned}$$



# EM algorithm - Jensen's Inequality

$$\begin{aligned}\log p(\mathbf{y}|\theta) &= \sum_{n=1}^N \ln \left( \sum_{\mathbf{z}} p(y_n, \mathbf{z}|\theta) \right) \\ &= \sum_{n=1}^N \ln \left( \sum_{\mathbf{z}} q(\mathbf{z}) \frac{p(y_n, \mathbf{z}|\theta)}{q(\mathbf{z})} \right) \\ &\geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \ln \left( \frac{p(y_n, \mathbf{z}|\theta)}{q(\mathbf{z})} \right)\end{aligned}$$



Assuming a value of  $\theta^{(t)}$ , what form  $q(\mathbf{z})$  should have to maximize

$$\ln p(\mathbf{y}|\theta^{(t)}) \geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \ln \left( \frac{p(\mathbf{z}|y_n, \theta^{(t)})p(y_n|\theta^{(t)})}{q(\mathbf{z})} \right)$$



Assuming a value of  $\theta^{(t)}$ , what form  $q(\mathbf{z})$  should have to maximize

$$\begin{aligned}\ln p(\mathbf{y}|\theta^{(t)}) &\geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \ln \left( \frac{p(\mathbf{z}|y_n, \theta^{(t)})p(y_n|\theta^{(t)})}{q(\mathbf{z})} \right) \\ &\geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \left[ \ln \left( \frac{p(\mathbf{z}|y_n, \theta^{(t)})}{q(\mathbf{z})} \right) + \ln p(y_n|\theta^{(t)}) \right]\end{aligned}$$



# Distribution $q(\mathbf{z})$

Assuming a value of  $\theta^{(t)}$ , what form  $q(\mathbf{z})$  should have to maximize

$$\begin{aligned}\ln p(\mathbf{y}|\theta^{(t)}) &\geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \ln \left( \frac{p(\mathbf{z}|y_n, \theta^{(t)})p(y_n|\theta^{(t)})}{q(\mathbf{z})} \right) \\ &\geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \left[ \ln \left( \frac{p(\mathbf{z}|y_n, \theta^{(t)})}{q(\mathbf{z})} \right) + \ln p(y_n|\theta^{(t)}) \right]\end{aligned}$$

organizing,

$$\sum_{n=1}^N \ln p(y_n|\theta^{(t)}) \geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \ln \left( \frac{p(\mathbf{z}|y_n, \theta^{(t)})}{q(\mathbf{z})} \right) + \sum_{n=1}^N \ln p(y_n|\theta^{(t)}),$$





# Distribution $q(\mathbf{z})$

Assuming a value of  $\theta^{(t)}$ , what form  $q(\mathbf{z})$  should have to maximize

$$\begin{aligned}\ln p(\mathbf{y}|\theta^{(t)}) &\geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \ln \left( \frac{p(\mathbf{z}|y_n, \theta^{(t)})p(y_n|\theta^{(t)})}{q(\mathbf{z})} \right) \\ &\geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \left[ \ln \left( \frac{p(\mathbf{z}|y_n, \theta^{(t)})}{q(\mathbf{z})} \right) + \ln p(y_n|\theta^{(t)}) \right]\end{aligned}$$

organizing,

$$\sum_{n=1}^N \ln p(y_n|\theta^{(t)}) \geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \ln \left( \frac{p(\mathbf{z}|y_n, \theta^{(t)})}{q(\mathbf{z})} \right) + \sum_{n=1}^N \ln p(y_n|\theta^{(t)}),$$

then,  $q(\mathbf{z}) = p(\mathbf{z}|y_n, \theta^{(t)})$ .



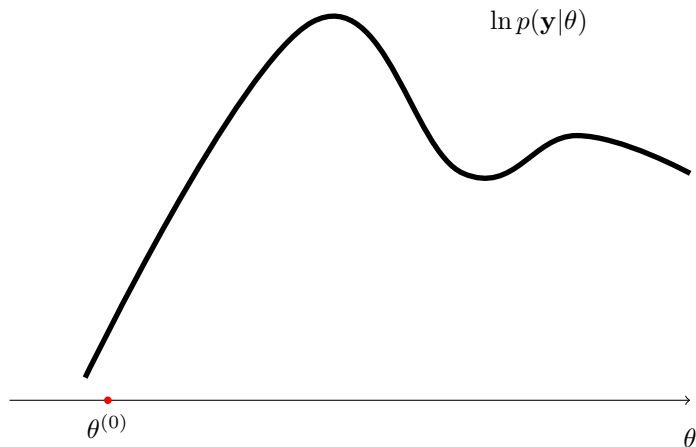
# Parameter estimation

If  $q(\mathbf{z}) = p(\mathbf{z}|y_n, \theta^{(t)})$ ,

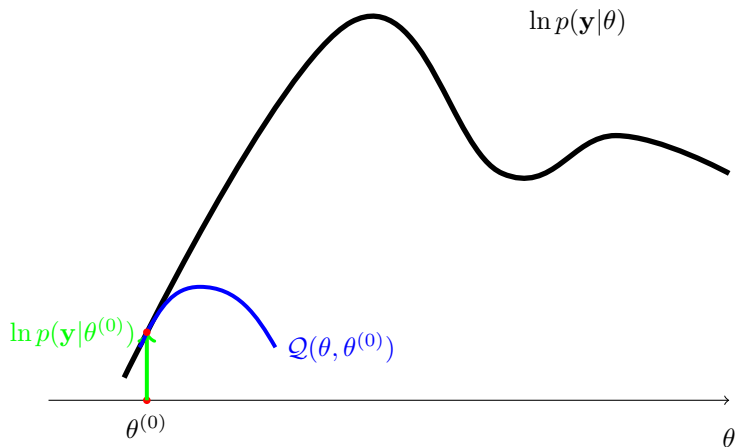
$$\begin{aligned}\theta^{(t+1)} &= \arg \max_{\theta} \sum_{n=1}^N \sum_{\mathbf{z}} p(\mathbf{z}|y_n, \theta^{(t)}) \ln \left( \frac{p(y_n, \mathbf{z}|\theta)}{p(\mathbf{z}|y_n, \theta^{(t)})} \right) \\ &= \arg \max_{\theta} \sum_{n=1}^N \mathbb{E}_{q(\mathbf{z})} [\ln (p(y_n, \mathbf{z}|\theta))] \\ &= \arg \max_{\theta} \mathcal{Q}(\theta, \theta^{(t)})\end{aligned}$$



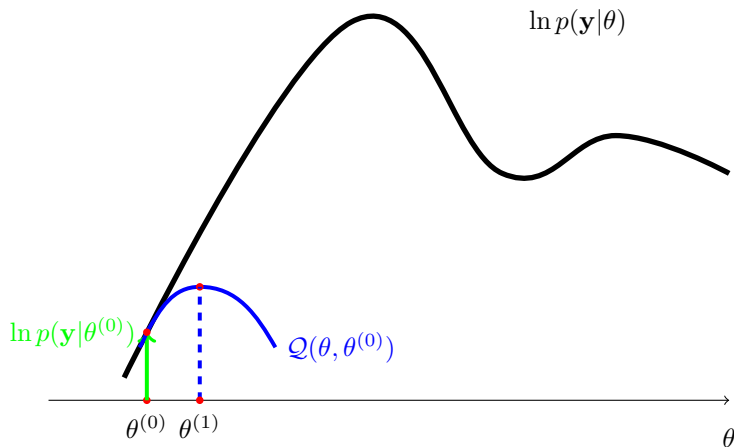
# EM algorithm - Visual explanation



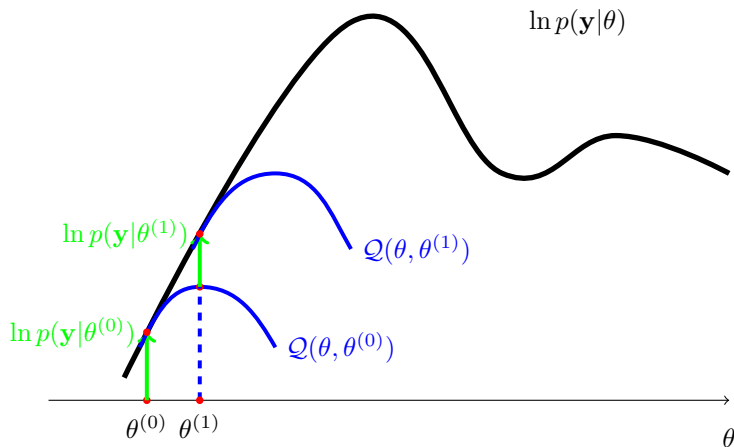
# EM algorithm - Visual explanation



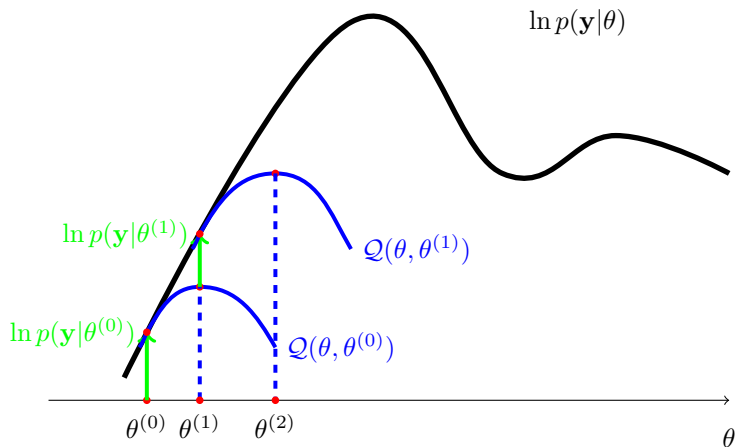
# EM algorithm - Visual explanation



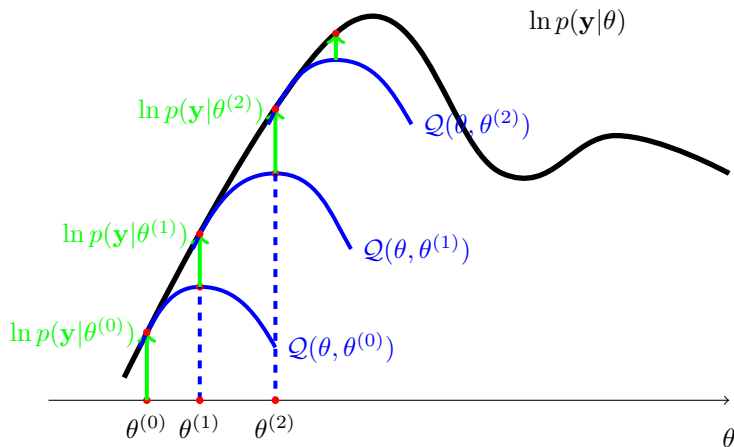
# EM algorithm - Visual explanation



# EM algorithm - Visual explanation



# EM algorithm - Visual explanation





# EM Algorithm - Pseudo-code

Given the joint distribution  $p(\mathbf{y}, \mathbf{z}|\theta)$ , the objective is to maximize  $p(\mathbf{y}|\theta)$  w.r.t.  $\theta$ .

**1** Select the initial parameters  $\theta^{(t)} \leftarrow \theta^{(0)}$ .

**2** Evaluate E-step,  $q(\mathbf{z}) = p(\mathbf{z}|\mathbf{y}, \theta^{(t)})$ .

**3** Evaluate M-step,

$$\theta^{(t+1)} = \arg \max_{\theta} \mathcal{Q}(\theta, \theta^{(t)})$$

**4** Verify convergence. If it doesn't satisfy, then

$$\theta^{(t)} \leftarrow \theta^{(t+1)},$$

return to step 2.



# EM algorithm - Remark

$$\log p(\mathbf{y}|\theta) \geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \ln \left( \frac{p(y_n, \mathbf{z}|\theta)}{q(\mathbf{z})} \right)$$

$$\sum_{n=1}^N \ln p(y_n|\theta) \geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \ln \left( \frac{p(\mathbf{z}|y_n, \theta)}{q(\mathbf{z})} \right) + \sum_{n=1}^N \ln p(y_n|\theta),$$

$$\sum_{n=1}^N \ln p(y_n|\theta) \geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \ln (p(y_n|\mathbf{z}, \theta)) + \sum_{\mathbf{z}} q(\mathbf{z}) \ln \left( \frac{p(\mathbf{z}|\theta)}{q(\mathbf{z})} \right).$$

$$\sum_{n=1}^N \ln p(y_n|\theta) \geq \sum_{n=1}^N \sum_{\mathbf{z}} q(\mathbf{z}) \ln (p(y_n, \mathbf{z}|\theta)) - \sum_{\mathbf{z}} q(\mathbf{z}) \ln (q(\mathbf{z})).$$



# EM algorithm - Summary

- 1 We just need to calculate the expected value of the joint distribution w.r.t. the posterior of the latent variables.
- 2 We are able to estimate the parameters  $\theta$  by maximizing an easier objective function.



1 EM algorithm

2 **Variational EM**

3 **Approximate inference for Neural Networks**

- Bayesian Deep Neural Networks
- Laplace Approximation



# Problem formulation

In this case,

$$p(\mathbf{z}|\mathbf{y}) = \frac{p(\mathbf{y}, \mathbf{z})}{p(\mathbf{y})}$$

is intractable. Most of the time because

$$\begin{aligned} p(\mathbf{y}) &= \int_{\mathbf{z}} p(\mathbf{y}, \mathbf{z}) \, d\mathbf{z} \\ &= \int_{\mathbf{z}} p(\mathbf{y}|\mathbf{z})p(\mathbf{z}) \, d\mathbf{z} \\ &= \mathbb{E}_{p(\mathbf{z})} [p(\mathbf{y}|\mathbf{z})] \end{aligned}$$

is difficult to calculate.



# Variational EM - Objective

We are interested in estimating  $p(\mathbf{z}|\mathbf{y})$ . This can be achieved by solving the following optimization problem,

$$q^*(\mathbf{z}) = \arg \min \text{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{y}))$$

where

$$\text{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{y})) = \int q(\mathbf{z}) \ln \frac{q(\mathbf{z})}{p(\mathbf{z}|\mathbf{y})} d\mathbf{z}$$



Analysing the KL divergence between  $q(\mathbf{z})$  and  $p(\mathbf{z}|\mathbf{y})$ ,

$$\begin{aligned}\text{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{y})) &= \int q(\mathbf{z}) \ln \frac{q(\mathbf{z})}{p(\mathbf{z}|\mathbf{y})} d\mathbf{z} \\ &= \int q(\mathbf{z}) \ln q(\mathbf{z}) d\mathbf{z} - \int q(\mathbf{z}) \ln p(\mathbf{z}|\mathbf{y}) d\mathbf{z} \\ &= \int q(\mathbf{z}) \ln q(\mathbf{z}) d\mathbf{z} - \int q(\mathbf{z}) \ln \frac{p(\mathbf{z}, \mathbf{y})}{p(\mathbf{y})} d\mathbf{z} \\ &= \int q(\mathbf{z}) \ln q(\mathbf{z}) d\mathbf{z} - \int q(\mathbf{z}) \ln p(\mathbf{z}, \mathbf{y}) d\mathbf{z} + \ln p(\mathbf{y})\end{aligned}$$



# Variational EM - The evidence lower bound

From previous equation we have,

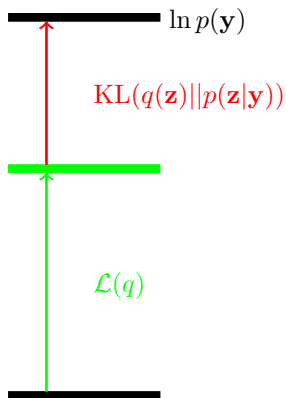
$$\begin{aligned}\ln p(\mathbf{y}) &= \text{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{y})) + \int q(\mathbf{z}) \ln p(\mathbf{z}, \mathbf{y}) d\mathbf{z} - \int q(\mathbf{z}) \ln q(\mathbf{z}) d\mathbf{z} \\ &= \text{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{y})) + \mathcal{L}(q).\end{aligned}$$

Given that  $\text{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{y})) \geq 0$ , then the lower bound of  $\ln p(\mathbf{y})$  is  $\mathcal{L}(q)$ .





# Variational EM - The evidence lower bound



# Variational EM - Summary

- 1 If the posterior  $p(\mathbf{z}|\mathbf{y})$  is intractable or complex (main issue is to calculate  $p(\mathbf{y})$ ).
- 2 We can approximate the posterior by minimizing  $\text{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{y}))$ , which results in maximizing  $\mathcal{L}(q)$ .
- 3 The E-step is an iterative process.



$$\text{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{y})) = \int q(\mathbf{z}) \ln \frac{q(\mathbf{z})}{p(\mathbf{z}|\mathbf{y})} \mathrm{d}\mathbf{z}$$

$$\text{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{y})) = \int q(\mathbf{z}) \ln q(\mathbf{z}) \mathrm{d}\mathbf{z} - \int q(\mathbf{z}) \ln p(\mathbf{z}, \mathbf{y}) \mathrm{d}\mathbf{z} + \ln p(\mathbf{y})$$

$$\text{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{y})) = - \int q(\mathbf{z}) \ln p(\mathbf{y}|\mathbf{z}) \mathrm{d}\mathbf{z} + \int q(\mathbf{z}) \ln \frac{q(\mathbf{z})}{p(\mathbf{z})} \mathrm{d}\mathbf{z} + \ln p(\mathbf{y})$$



1 EM algorithm

2 Variational EM

3 **Approximate inference for Neural Networks**

- Bayesian Deep Neural Networks
- Laplace Approximation



## 1 EM algorithm

## 2 Variational EM

## 3 Approximate inference for Neural Networks

- Bayesian Deep Neural Networks
- Laplace Approximation



# Approx. NNs - Problem definition

Let's assume we have a dataset  $\mathcal{D} = \{\mathbf{X}, \mathbf{Y}\}$ , with  $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$  and  $\mathbf{Y} = \{\mathbf{y}_1, \dots, \mathbf{y}_N\}$

$$\begin{aligned}\text{KL}(q_{\theta}(\boldsymbol{\omega}) \| p(\boldsymbol{\omega} | \mathcal{D})) &\propto - \int q_{\theta}(\boldsymbol{\omega}) \log p(\mathbf{Y} | \mathbf{X}, \boldsymbol{\omega}) d\boldsymbol{\omega} + \text{KL}(q_{\theta}(\boldsymbol{\omega}) \| p(\boldsymbol{\omega})), \\ &= - \sum_{i=1}^N \int q_{\theta}(\boldsymbol{\omega}) \log p(\mathbf{y}_i | \mathbf{f}^{\boldsymbol{\omega}}(\mathbf{x}_i)) d\boldsymbol{\omega} + \text{KL}(q_{\theta}(\boldsymbol{\omega}) \| p(\boldsymbol{\omega})),\end{aligned}$$

where  $\boldsymbol{\omega}$  is a vector composed by the stacking of all weights.



# Approx. NNs - Problem definition

- 1 The summed-over terms log-likelihood are not tractable for BNNs with more than a single hidden layer.
- 2 This objective requires us to perform computations over the entire dataset, which can be too costly for large  $N$ .



We can reduce the data size problem by data sub-sampling (also referred to as mini-batch optimization).

$$\hat{\mathcal{L}}_{\text{VI}}(\theta) := -\frac{N}{M} \sum_{i \in S} \int q_{\theta}(\omega) \log p(\mathbf{y}_i | \mathbf{f}^{\omega}(\mathbf{x}_i)) d\omega + \text{KL}(q_{\theta}(\omega) \| p(\omega))$$

with a random index set  $S$  of size  $M$ .





# Approx. NNs - Expected value solution

We can reduce the data size problem by data sub-sampling (also referred to as mini-batch optimization).

$$\mathbb{E}_{q_{\theta}(\boldsymbol{\omega})}[\log p(\mathbf{y}_i | \mathbf{f}^{\boldsymbol{\omega}}(\mathbf{x}_i))] = \int q_{\theta}(\boldsymbol{\omega}) \log p(\mathbf{y}_i | \mathbf{f}^{\boldsymbol{\omega}}(\mathbf{x}_i)) d\boldsymbol{\omega}$$

$$q_{\theta}(\boldsymbol{\omega}) = \prod_{i=1}^L q_{\theta}(\mathbf{W}_i) = \prod_{i=1}^L \prod_{j=1}^{K_i} \prod_{k=1}^{K_{i+1}} q(w_{ijk}) = \prod_{i,j,k} \mathcal{N}(w_{ijk}; m_{ijk}, \sigma_{ijk}^2)$$

$$\mathbb{E}_{q_{\theta}(\boldsymbol{\omega})}[\log p(\mathbf{y}_i | \mathbf{f}^{\boldsymbol{\omega}}(\mathbf{x}_i))] \approx \frac{1}{M} \sum_{m=1}^M q_{\theta}(\boldsymbol{\omega}_j) \log p(\mathbf{y}_i | \mathbf{f}^{\boldsymbol{\omega}}(\mathbf{x}_i))$$



# Approx. NNs - Expected value solution

A better way is to use Pathwise Gradient Estimator

$$\mathbb{E}_{q_{\theta}(\omega)} [f_{\theta}(\omega)] \longrightarrow \mathbb{E}_{q_0(\epsilon)} [f_{\theta}(g(\theta; \epsilon))]$$

$$q(w_{ijk}) = g(\theta_{ijk}, \epsilon_{ijk}) = m_{ijk} + \sigma_{ijk}\epsilon_{ijk}$$

then

$$\hat{\mathcal{L}}_{\text{VI}}(\theta) := -\frac{N}{M} \sum_{i \in S} \int q_{\theta}(\omega) \log p(\mathbf{y}_i | \mathbf{f}^{\omega}(\mathbf{x}_i)) d\omega + \text{KL}(q_{\theta}(\omega) \| p(\omega))$$

becomes

$$\hat{\mathcal{L}}_{\text{MC}}(\theta) = -\frac{N}{M} \sum_{i \in S} \log p(\mathbf{y}_i | \mathbf{f}^{g(\theta, \epsilon)}(\mathbf{x}_i)) + \text{KL}(q_{\theta}(\omega) \| p(\omega))$$

where  $\mathbb{E}_{S, \epsilon} [\hat{\mathcal{L}}_{\text{MC}}(\theta)] = \mathcal{L}_{\text{VI}}(\theta)$ .



---

**Algorithm 1** Minimise divergence between  $q_{\theta}(\omega)$  and  $p(\omega|X, Y)$ 

---

- 1: Given dataset  $\mathbf{X}, \mathbf{Y}$ ,
- 2: Define learning rate schedule  $\eta$ ,
- 3: Initialise parameters  $\theta$  randomly.
- 4: **repeat**
- 5:   Sample  $M$  random variables  $\hat{\epsilon}_i \sim p(\epsilon)$ ,  $S$  a random subset of  $\{1, \dots, N\}$  of size  $M$ .
- 6:   Calculate stochastic derivative estimator w.r.t.  $\theta$ :

$$\widehat{\Delta\theta} \leftarrow -\frac{N}{M} \sum_{i \in S} \frac{\partial}{\partial \theta} \log p(\mathbf{y}_i | \mathbf{f}^{q(\theta, \hat{\epsilon}_i)}(\mathbf{x}_i)) + \frac{\partial}{\partial \theta} \text{KL}(q_{\theta}(\omega) || p(\omega)).$$

- 7:   Update  $\theta$ :  
       $\theta \leftarrow \theta + \eta \widehat{\Delta\theta}$ .
  - 8: **until**  $\theta$  has converged.
- 



$$\begin{aligned}\hat{\mathbf{y}} &= \hat{\mathbf{h}}\mathbf{M}_2 \\ &= (\mathbf{h} \odot \hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2 \\ &= (\mathbf{h} \cdot \text{diag}(\hat{\boldsymbol{\epsilon}}_2))\mathbf{M}_2 \\ &= \mathbf{h}(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2) \\ &= \sigma(\hat{\mathbf{x}}\mathbf{M}_1 + \mathbf{b})(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2) \\ &= \sigma((\mathbf{x} \odot \hat{\boldsymbol{\epsilon}}_1)\mathbf{M}_1 + \mathbf{b})(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2) \\ &= \sigma(\mathbf{x}(\text{diag}(\hat{\boldsymbol{\epsilon}}_1)\mathbf{M}_1) + \mathbf{b})(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2)\end{aligned}$$



## 1 EM algorithm

## 2 Variational EM

## 3 Approximate inference for Neural Networks

- Bayesian Deep Neural Networks

- Laplace Approximation



# Approx. NNs - Laplace approximation

The standard Laplace approximation is obtained by taking the second-order Taylor expansion around a mode of a distribution.

If we approximate the log posterior over the weights of a network given some data  $\mathcal{D}$  around a MAP estimate  $\omega^*$ , we obtain

$$\log p(\omega|\mathcal{D}) \approx \log p(\omega^*|\mathcal{D}) - \frac{1}{2} (\omega - \omega^*)^\top \overline{H} (\omega - \omega^*),$$

where  $\overline{H} = \mathbb{E}[H]$  is the average Hessian of the negative log posterior. The posterior over the weights is then approximated as Gaussian:

$$\omega \sim \mathcal{N}(\omega^*, \overline{H}^{-1})$$





Bishop, C. M. (2006).

*Pattern Recognition and Machine Learning (Information Science and Statistics).*

Springer-Verlag, Berlin, Heidelberg.



Gal, Y. (2016).

*Uncertainty in Deep Learning.*

PhD thesis, University of Cambridge.



Ritter, H., Botev, A., and Barber, D. (2018).

A Scalable Laplace Approximation for Neural Networks.

*In International Conference on Learning Representations.*